

National University of Computer & Emerging Sciences

CS 3001 - COMPUTER NETWORKS

Lecture 13

Chapter 3

11th March, 2025

Nauman Moazzam Hayat

nauman.moazzam@lhr.nu.edu.pk

Office Hours: 11:30 am till 01:00 pm (Every Tuesday & Thursday)

Chapter 3

Transport Layer

A note on the use of these PowerPoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2023
J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top-Down Approach

8th edition

Jim Kurose, Keith Ross
Pearson, 2020

Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- **Connection-oriented transport: TCP**
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- Principles of congestion control
- TCP congestion control



TCP Sender (simplified)

event: data received from application

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
 - think of timer as for oldest unACKed segment
 - expiration interval: **TimeoutInterval**

event: timeout

- retransmit segment that caused timeout
- restart timer

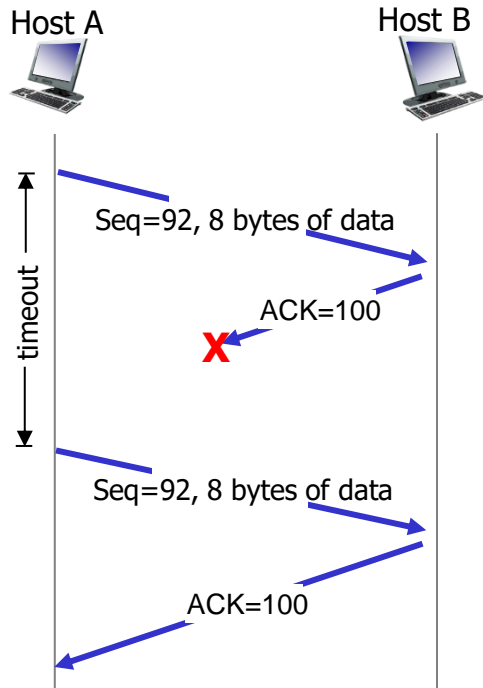
event: ACK received (*cumulative ACK*)

- if ACK acknowledges previously unACKed segments
 - update what is known to be ACKed
 - start timer if there are still unACKed segments

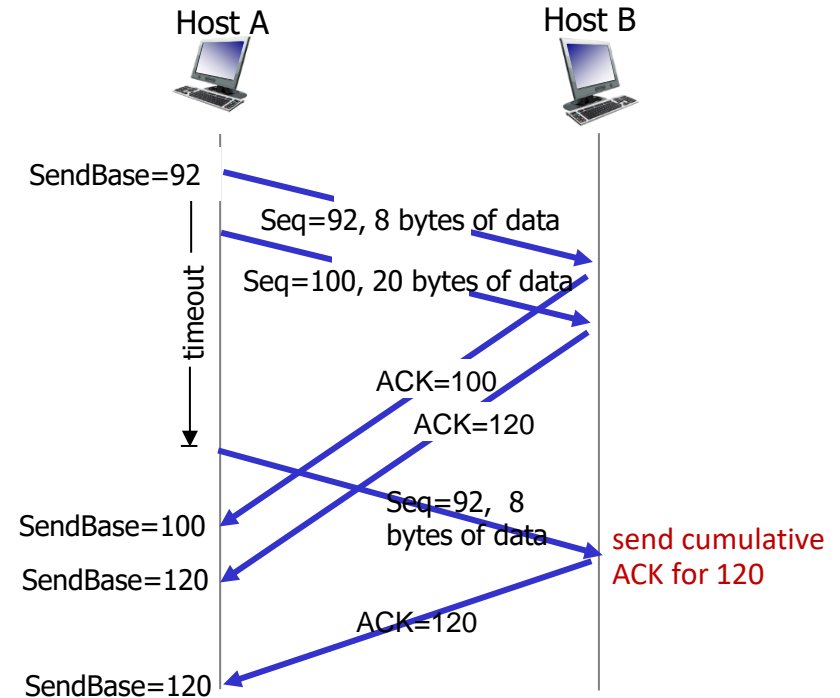
TCP Receiver: ACK generation [RFC 5681]

<i>Event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

TCP: retransmission scenarios

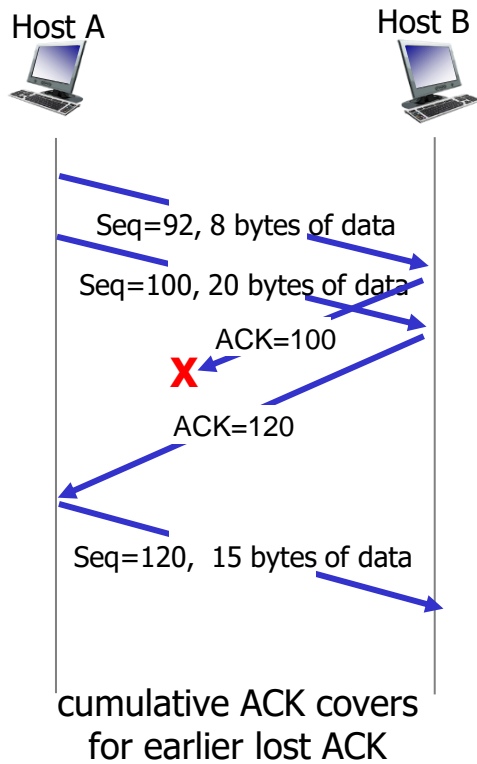


lost ACK scenario



premature timeout

TCP: retransmission scenarios



TCP fast retransmit


TCP fast retransmit

if sender receives 3 additional ACKs for same data (“triple duplicate ACKs”), resend unACKed segment with smallest seq #

- likely that unACKed segment lost, so don't wait for timeout



Receipt of three duplicate ACKs indicates 3 segments received after a missing segment – lost segment is likely. So retransmit!



So is TCP Go Back-N or Selective Repeat???


- *TCP (error / loss recovery mechanism) is best categorized as a **hybrid** of GBN & SR*
- *Please refer to Page 246 of the Course Textbook for further details.*

Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- **Connection-oriented transport: TCP**
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- Principles of congestion control
- TCP congestion control

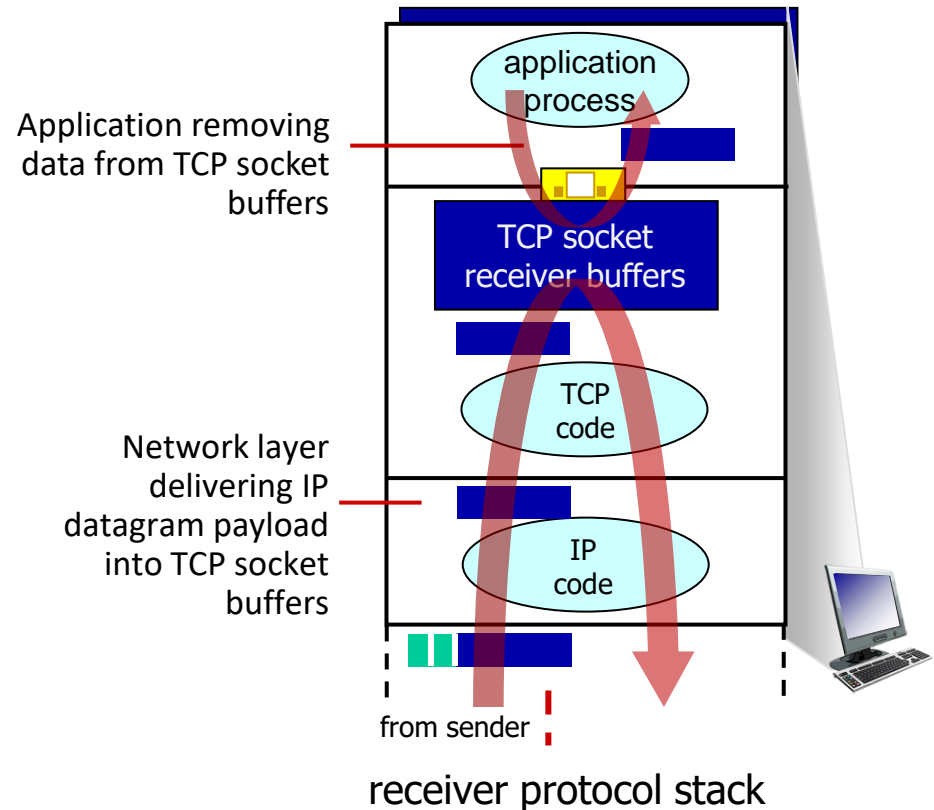


Flow Control vs Congestion Control (Courtesy Rutgers University)



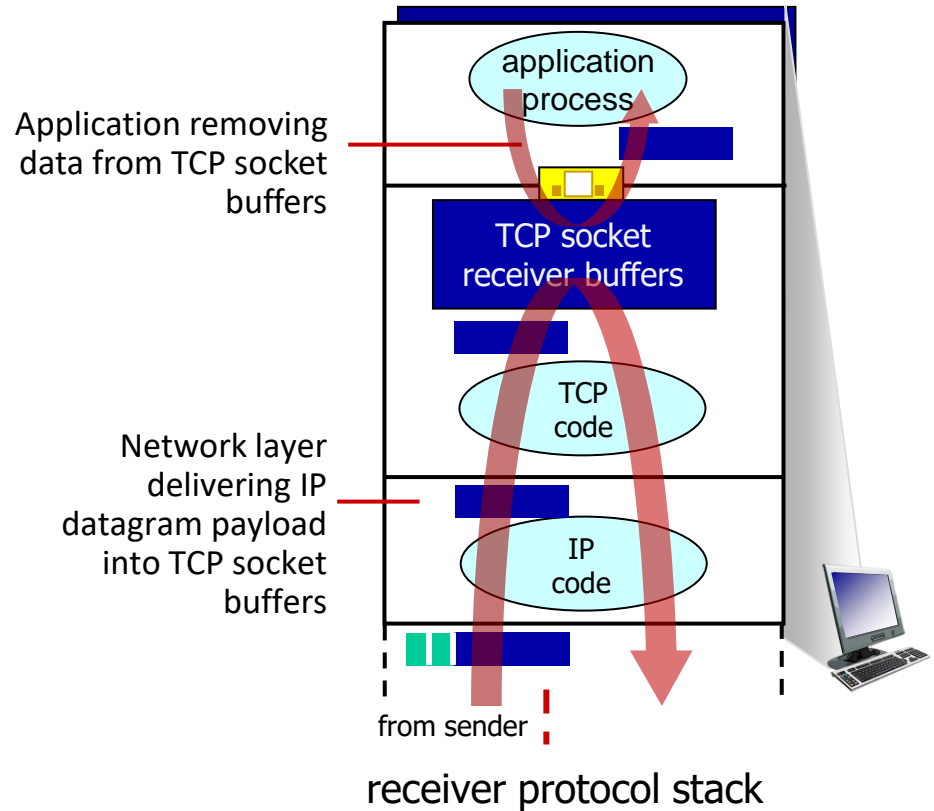
TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



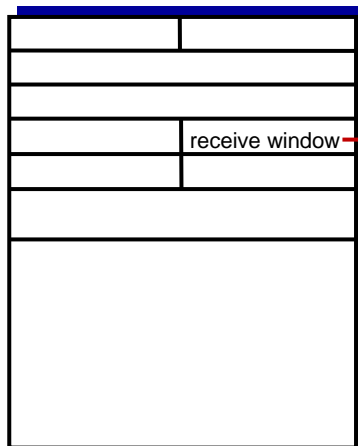
TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



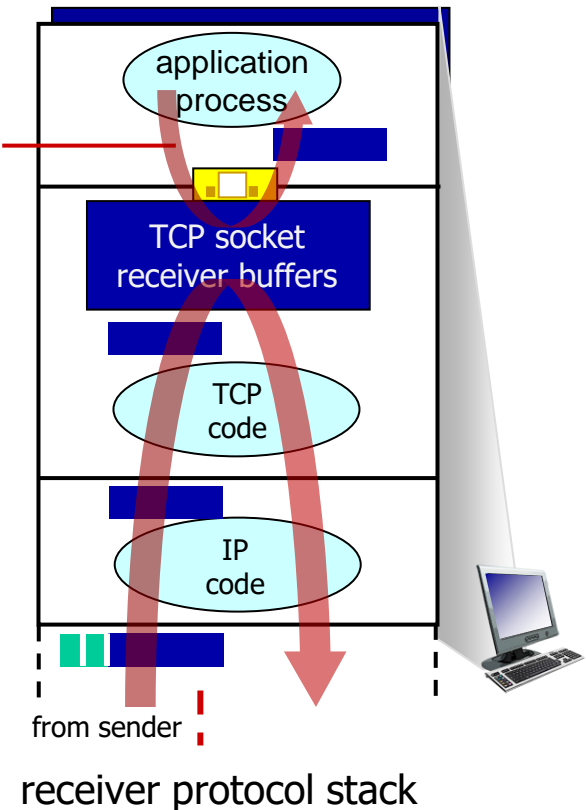
TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



flow control: # bytes receiver willing to accept

Application removing data from TCP socket buffers

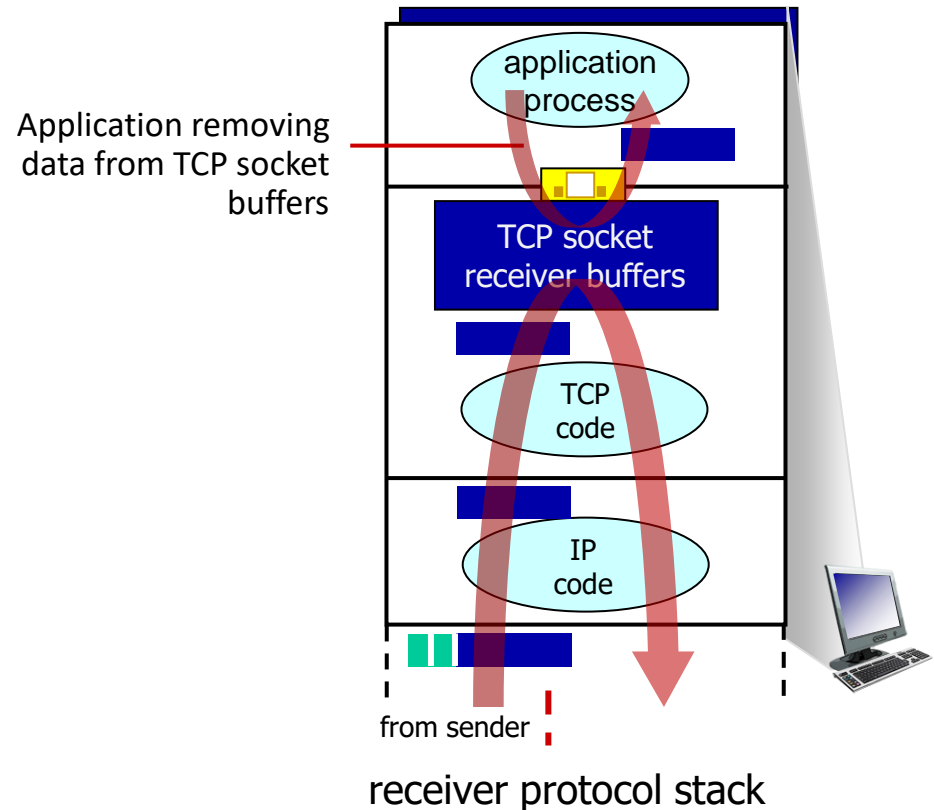


TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?

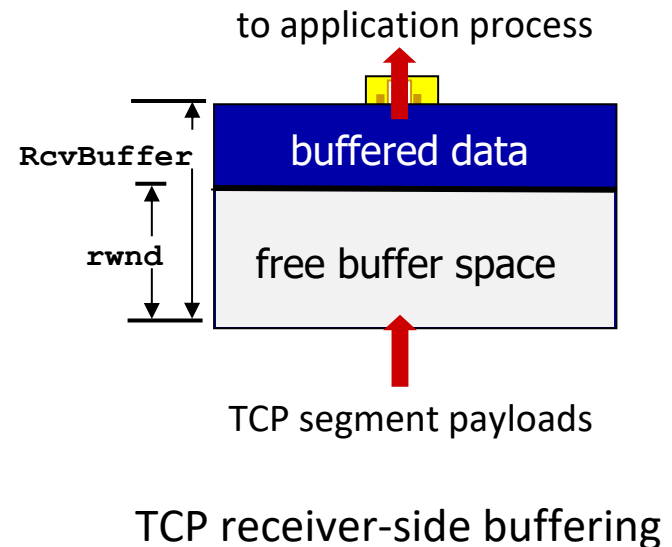
flow control

receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast



TCP flow control


- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
 - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
 - many operating systems auto-adjust **RcvBuffer**
- sender limits amount of unACKed (“in-flight”) data to received **rwnd**
- guarantees receive buffer will not overflow



TCP flow control

- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
 - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
 - many operating systems auto-adjust **RcvBuffer** (**i.e. dynamic**)
- sender limits amount of unACKed (“in-flight”) data to received **rwnd**
- guarantees receive buffer will not overflow
- (**Assumption: TCP receiver discards out-of-order segments**)

flow control: # bytes receiver willing to accept



TCP segment format

TCP Flow Control - Implementation

- ❖ TCP provides flow control by having the sender maintain a dynamic variable called the receive window (**rwnd**)

(Since TCP is not permitted to overflow the allocated buffer, **we must have**)

$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{RcvBuffer}$ (where:

- **RcvBuffer** is the size of the buffer allocated at the receiver for this connection,
- **LastByteRcvd** is the number of the last byte in the data stream received from the network and placed in the receive buffer, &
- **LastByteRead** is the number of the last byte in the data stream read by the application process from the receive buffer), **thus**

$$\text{rwnd} = \text{RcvBuffer} - [\text{LastByteRcvd} - \text{LastByteRead}]$$

- ❖ While the receiver is keeping track of many variables as seen above, the sender is keeping track of primarily two variables, (i.e. **LastByteSent** & **LastByteAcked**)

$\text{LastByteSent} - \text{LastByteAcked} \leq \text{rwnd}$ (i.e. the **UnAckedData**.)

- By maintaining this throughout the connection's life, i.e. keeping the **UnAckedData** less than or equal to the **rwnd**, the sender ensures it doesn't overflow the buffer at the receiver

TCP Flow Control - Issue

Issue: One minor technical problem with this scheme.

- To see this, suppose Host B's (receiver) receive buffer becomes full so that $rwnd = 0$.
- After advertising $rwnd = 0$ to Host A (sender), also suppose that B has nothing to send to A.
- As the application process at B empties the buffer, TCP does not send new segments with new $rwnd$ values to Host A (indeed, TCP sends a segment to Host A only if it has data to send or if it has an acknowledgment to send)

Therefore, Host A is never informed that some space has opened up in Host B's receive buffer (i.e. **Host A is blocked and can transmit no more data!**)

- To **solve** this problem, the TCP specification requires Host A to continue to send segments with one data byte when B's receive window is zero.
- These segments will be acknowledged by the receiver.
- Eventually the buffer will begin to empty and the acknowledgments will contain a non-zero $rwnd$ value.

Assignment # 3 (Chapter - 3)

- *3rd Assignment will be uploaded on Google Classroom on Thursday, 13th March, 2025, in the Stream - Announcement Section*
- *Due Date: ~~Thursday, 20th March~~ Tuesday, 25th March, 2025
(Handwritten solutions to be submitted during the lecture; deadline extended due to LAB midterms next week)*
- *Please read **all the instructions** carefully in the uploaded Assignment document, follow & submit accordingly*

Quiz # 3 (Chapter - 3)

- *On: ~~Thursday, 20th March, 2025~~ , Tuesday, 25th March, 2025 (During the lecture; deadline extended due to LAB midterms next week)*
- *Quiz to be taken during own section class only*